

MILESTONE 3

From insight to solution

Increasing voice input adoption for ChatGPT

Students and job seekers, 18–27, Android, mobile-first — India

Simran · Cohort 52 · 2-week sprint

Voice is available. People still type.

Students and job seekers use ChatGPT daily for effortful work and type almost all of it.

54.8%

almost always type

9.7%

mostly use voice

0%

always use voice

Survey, n=31 · Milestone 2

The insight

Voice is not unavailable — it does not feel safe, accurate, or controllable in the shared, on-the-go places where this segment actually uses their phone.

Three blockers, in the users' own words

1

Social self-consciousness

Speaking out loud near classmates, family, or on a crowded metro feels exposing. The biggest single blocker.

2

Low trust in accuracy

32.3% said it “may misunderstand my voice.” Hinglish and code-mixed speech are where it breaks.

3

No live feedback

You cannot see what was heard until it is already sent. One respondent asked directly for text while speaking.

Each of the three solution directions on the next slide attacks one of these.

SECTION 1

Three solution directions

Each attacks a different root cause. All three are viable.



Live transcript + edit-before-send

A real-time, editable caption as you speak. Low-confidence words are flagged; tap to fix or switch a language chip before anything sends.

Targets: Accuracy distrust · no live feedback



Ambient / home-screen entry

A lock-screen or home-screen voice shortcut, the way Gemini and Assistant do, so voice does not require already being inside a chat.

Targets: Discoverability gap



Regional-accent confidence mode

Pick a preferred accent or dialect at onboarding, with a visible “tuned for Hinglish” badge to build trust before you ever speak.

Targets: Distrust of code-mixed recognition

A is the chosen direction — the next slide shows why.

Trade-offs: impact, effort, confidence

Direction	Impact	Effort	Confidence	Reasoning
A — Live transcript + edit-before-send	High	Medium	High	Hits the two highest-scoring blockers and was explicitly requested by a respondent. Ships inside the existing chat surface.
B — Ambient / home-screen entry	Med–High	High	Medium	Would lift discovery, but a user who distrusts accuracy will not keep using an easier door into the same problem.
C — Regional-accent confidence mode	Medium	Med–High	Medium	Depends on model investment we cannot scope yet, and the trust claim is unverifiable until recognition actually improves.

Chosen: Direction A. No new OS surface, no accent-model investment, and it lives inside the chat people already open every day.

Why Direction A goes first

What it gets right

- Answers the blockers users actually named
- No new OS-level surface — unlike Direction B
- No accent-model investment — unlike Direction C
- Shortest path from “typing is the default” to a voice flow people repeat

What we are not building yet

- No home-screen or lock-screen surface — that is B
- No change to the speech model itself
- No voice output — the reply stays text, which is what a user in a shared space wants

Where Direction A is honest about its limits

Self-consciousness was the biggest blocker, and a transcript does not make a crowded metro quieter. What it removes is the consequence of speaking softly: nothing sends until you have read it back, so a half-heard sentence is a two-tap fix, not a message you apologise for. That is a partial answer — a whisper-optimised mode is the follow-up if the metrics say self-consciousness still dominates.

What would change our mind: if time-to-send lands worse than typing, or the correction rate stays flat after four weeks, Direction C moves ahead.


SECTION 3

User flow — discovery to use

Voice still starts at the mic icon in the message bar. Everything new happens between tapping it and sending.

User Flow — Discovering & Using Voice

Students & job seekers, 18-27, Android mobile-first



SECTION 3

The flow has to survive real conditions

A happy path is not a design. These are the states the feature fails into.



Mic permission denied

Fall back to the keyboard silently. Re-offer after three sessions, never as a blocking prompt.



User goes quiet mid-sentence

Keep listening 3 seconds, then move to review with whatever was captured. Never discard audio.



Noisy room, most words flagged

Show the transcript anyway, with a prompt to switch language or record again. Flagging beats silently guessing.



User taps away mid-recording


Save the draft into the message bar so returning does not lose the attempt.

SECTION 4

Annotated wireframes

The mic does not move. What changes is everything after the tap.

Screen 1 — Entry point



Screen 1 — Entry Point

Discovery: same tap-to-start users already know

- 1 **No new gesture to learn.** The mic icon stays exactly where it is today, inside the message bar — this keeps the existing entry point rather than adding a new ambient surface (that is a separate, larger effort — see Direction B in trade-offs).
- The change starts **after** the tap: what happens next (Screens 2–4) is what fixes the trust and feedback gap, without asking users to relearn where voice lives.

Screen 2 — Live transcript



Screen 2 — Live Transcript (Listening)

Root causes addressed: no live feedback · trust in accuracy

- 2 Captions appear **as the user speaks**, in the language/script actually used (Hindi/Hinglish/English mixed) — closing the "I have no idea what it heard" gap called out directly in user research.
- 3 Words the model is **less confident about** are underlined live. This previews Screen 3's edit step and sets expectation that correction is normal, not a failure state.
- No audio is sent until the user confirms (Screen 3) — so a user speaking quietly or in a shared room is never "caught out" by an accidental send.

SECTION 4

Annotated wireframes (continued)

Correction is one tap, and nothing is sent until the user says so.

Screen 3 — Review and edit




Screen 3 — Review & Edit

Root causes addressed: accuracy trust · habit & control · typing as fallback

- 4 The transcript is a normal **editable text field**, not a locked result — tapping a flagged word opens the keyboard right there, so a misheard word is a small fix instead of a scrap-and-retype.
- 5 A **language chip** (English / Hindi / Hinglish) lets the user correct the recognition mode in one tap when code-mixed speech is misread as pure English or vice-versa.
- 6 Nothing is sent until the user taps **Send** — "Redo" re-records instead of forcing a full restart from the chat screen, so control stays with the user throughout.

Screen 4 — Sent and answered



Screen 4 — Sent & Answered

Closing the loop, quietly building the habit

- 7 A small **"voice" tag** marks the message as sent via voice — an unobtrusive, private cue that reinforces to the user that voice "worked", without announcing it to anyone nearby.
- 8 The reply is indistinguishable from a typed turn — voice is a second way to compose the same message, not a separate mode with different capabilities.
- Behind the scenes, this turn logs whether the transcript was edited, how many words were corrected, and the language chip used — the inputs for the success metrics on page 3.

Success metrics

North star: voice send-through rate — the share of voice attempts that end in a sent message. The trust signal underneath it is corrections per message, which should fall week over week as users stop bracing for errors.

Stage	Metric	Baseline	90-day target
Discoverability	% of chat sessions where the mic is tapped	~10%	20%
Activation	% of voice attempts that reach Send	Not tracked	75%
Trust	Corrections per sent voice message	New metric	Falling 4 weeks running
Engagement	Voice messages per active voice user per week	Not tracked	3+
Retention	Week-4 retention, voice vs text-only	Not tracked	Voice ≥ text-only
Segment health	Send-through, Hindi/Hinglish vs English	Not tracked	Gap under 10 pts

Baselines are from the Milestone 2 survey (n=31) and are directional only — they would be re-established from product analytics before the test runs.

Guardrails and risks

What would tell us this is making things worse, and what we are not sure about.

Guardrails

- Caption latency under ~300ms per word — a laggy caption breaks the whole trust effect
- Time-to-send for voice does not exceed typing the same message
- No rise in overall session abandonment
- No drop in messages sent per session — voice should add, not cannibalise

Risks and assumptions

- The review step may read as friction, not safety — watch drop-off between review and Send
- Confidence scores may not match real errors — sample 200 transcripts against human ground truth
- Self-consciousness may dominate regardless — interview users who open voice and abandon
- Baselines come from n=31, not analytics — re-establish before testing

HOW WE WOULD TEST IT

Four weeks, 50/50, Android India

Audience	Android users in India, 18–27
Split	Control keeps today's voice input; treatment gets live transcript + review
Duration	Four weeks — long enough for correction rate to trend, not just spike
Primary metric	Voice send-through rate
Rollback	Automatic if session abandonment rises

If it works, Directions B and C are the follow-ups — in that order.